

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split, GridSearchCV, RandomizedSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
```

```
!pip install shap
```

```
Collecting shap
  Downloading shap-0.46.0-cp310-cp310-manylinux_2_12_x86_64.manylinux2010_x86_64.manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata
Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages (from shap) (1.26.4)
Requirement already satisfied: scipy in /usr/local/lib/python3.10/dist-packages (from shap) (1.13.1)
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.10/dist-packages (from shap) (1.3.2)
Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (from shap) (2.1.4)
Requirement already satisfied: tqdm>=4.27.0 in /usr/local/lib/python3.10/dist-packages (from shap) (4.66.5)
Requirement already satisfied: packaging>20.9 in /usr/local/lib/python3.10/dist-packages (from shap) (24.1)
Collecting slicer==0.0.8 (from shap)
  Downloading slicer-0.0.8-py3-none-any.whl.metadata (4.0 kB)
Requirement already satisfied: numba in /usr/local/lib/python3.10/dist-packages (from shap) (0.60.0)
Requirement already satisfied: cloudpickle in /usr/local/lib/python3.10/dist-packages (from shap) (2.2.1)
Requirement already satisfied: llvmlite<0.44,>=0.43.0dev0 in /usr/local/lib/python3.10/dist-packages (from numba->shap) (0.43.0)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.10/dist-packages (from pandas->shap) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas->shap) (2024.1)
Requirement already satisfied: tzdata>=2022.1 in /usr/local/lib/python3.10/dist-packages (from pandas->shap) (2024.1)
Requirement already satisfied: joblib>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from scikit-learn->shap) (1.4.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn->shap) (3.5.0)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.8.2->pandas->shap) (1.16.0)
Downloading shap-0.46.0-cp310-cp310-manylinux_2_12_x86_64.manylinux2010_x86_64.manylinux_2_17_x86_64.manylinux2014_x86_64.whl (540 kB)
540.1/540.1 kB 3.7 MB/s eta 0:00:00
Downloading slicer-0.0.8-py3-none-any.whl (15 kB)
Installing collected packages: slicer, shap
Successfully installed shap-0.46.0 slicer-0.0.8
```

```
import shap
```

```
df = pd.read_csv("/content/Churn_Modelling.csv")
print(df.head())
print(df.info())
print(df.describe())
```

```
EstimatedSalary  Exited
0      101348.88      1
1      112542.58      0
2      113931.57      1
3       93826.63      0
4       79084.10      0
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
 #   Column              Non-Null Count  Dtype
---  -
 0   RowNumber           10000 non-null  int64
 1   CustomerId          10000 non-null  int64
 2   Surname              10000 non-null  object
 3   CreditScore          10000 non-null  int64
 4   Geography            10000 non-null  object
 5   Gender               10000 non-null  object
 6   Age                  10000 non-null  int64
 7   Tenure               10000 non-null  int64
```

| | | | | | |
|-----|-------------|--------------|------------|-----------|-----------|
| 25% | 2500.75000 | 1.562853e+07 | 584.000000 | 32.000000 | 3.000000 |
| 50% | 5000.50000 | 1.569074e+07 | 652.000000 | 37.000000 | 5.000000 |
| 75% | 7500.25000 | 1.575323e+07 | 718.000000 | 44.000000 | 7.000000 |
| max | 10000.00000 | 1.581569e+07 | 850.000000 | 92.000000 | 10.000000 |

| | Balance | NumOfProducts | HasCrCard | IsActiveMember | \ |
|-------|---------------|---------------|--------------|----------------|---|
| count | 10000.000000 | 10000.000000 | 10000.000000 | 10000.000000 | |
| mean | 76485.889288 | 1.530200 | 0.70550 | 0.515100 | |
| std | 62397.405202 | 0.581654 | 0.45584 | 0.499797 | |
| min | 0.000000 | 1.000000 | 0.00000 | 0.000000 | |
| 25% | 0.000000 | 1.000000 | 0.00000 | 0.000000 | |
| 50% | 97198.540000 | 1.000000 | 1.00000 | 1.000000 | |
| 75% | 127644.240000 | 2.000000 | 1.00000 | 1.000000 | |
| max | 250898.090000 | 4.000000 | 1.00000 | 1.000000 | |

| | EstimatedSalary | Exited |
|-------|-----------------|--------------|
| count | 10000.000000 | 10000.000000 |
| mean | 100090.239881 | 0.203700 |
| std | 57510.492818 | 0.402769 |
| min | 11.580000 | 0.000000 |
| 25% | 51002.110000 | 0.000000 |
| 50% | 100193.915000 | 0.000000 |
| 75% | 149388.247500 | 0.000000 |
| max | 199992.480000 | 1.000000 |

```
df = df.drop(['RowNumber', 'CustomerId', 'Surname'], axis=1)
df = pd.get_dummies(df, drop_first=True)
X = df.drop('Exited', axis=1)
y = df['Exited']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

```
# Train a Random Forest Classifier
rf_model = RandomForestClassifier(random_state=42)
rf_model.fit(X_train, y_train)
y_pred = rf_model.predict(X_test)
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))

print("\nClassification Report:")
print(classification_report(y_test, y_pred))

print("\nAccuracy Score:")
print(accuracy_score(y_test, y_pred))
```

Confusion Matrix:

```
[[2328  88]
 [ 310 274]]
```

Classification Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.88 | 0.96 | 0.92 | 2416 |
| 1 | 0.76 | 0.47 | 0.58 | 584 |
| accuracy | | | 0.87 | 3000 |
| macro avg | 0.82 | 0.72 | 0.75 | 3000 |
| weighted avg | 0.86 | 0.87 | 0.85 | 3000 |

Accuracy Score:
0.8673333333333333

```
# Define the parameter grid for RandomSearch
param_grid = {
    'n_estimators': [50, 100, 200],
    'max_features': ['auto', 'sqrt'],
    'max_depth': [10, 20, 30, None],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'bootstrap': [True, False]
}

# RandomizedSearchCV
rf_random = RandomizedSearchCV(estimator=rf_model, param_distributions=param_grid, n_iter=100, cv=3, verbose=2, random_state=42, n_jobs=-1)
```

```

rf_random.fit(X_train, y_train)

print(f"Best parameters found by RandomizedSearch: {rf_random.best_params_}")
best_rf_model = rf_random.best_estimator_

y_pred_best = best_rf_model.predict(X_test)

print("Confusion Matrix (Best Model):")
print(confusion_matrix(y_test, y_pred_best))

print("\nClassification Report (Best Model):")
print(classification_report(y_test, y_pred_best))

print("\nAccuracy Score (Best Model):")
print(accuracy_score(y_test, y_pred_best))

estimator._validate_params()
File "/usr/local/lib/python3.10/dist-packages/sklearn/base.py", line 638, in _validate_params
    validate_parameter_constraints(
File "/usr/local/lib/python3.10/dist-packages/sklearn/utils/_param_validation.py", line 96, in validate_parameter_constraints
    raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The 'max_features' parameter of RandomForestClassifier must be an int in the r

-----
80 fits failed with the following error:
Traceback (most recent call last):
  File "/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py", line 729, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "/usr/local/lib/python3.10/dist-packages/sklearn/base.py", line 1145, in wrapper
    estimator._validate_params()
  File "/usr/local/lib/python3.10/dist-packages/sklearn/base.py", line 638, in _validate_params
    validate_parameter_constraints(
  File "/usr/local/lib/python3.10/dist-packages/sklearn/utils/_param_validation.py", line 96, in validate_parameter_constraints
    raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The 'max_features' parameter of RandomForestClassifier must be an int in the r

warnings.warn(some_fits_failed_message, FitFailedWarning)
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_search.py:979: UserWarning: One or more of the test scores are non-f
0.85842867 0.86314308 0.86157155 0.86342859 0.86485724 nan
0.86357134 nan 0.85900018 nan nan nan
nan nan nan nan 0.86100016 nan
nan nan 0.86314308 nan 0.86400004 0.86057183
nan nan 0.86028559 0.86199993 0.86199975 nan
0.86342853 0.86157142 0.86285708 0.86342853 0.85614287 0.85585712
nan nan 0.86214299 nan 0.85700014 nan
0.86314308 nan nan nan 0.8574284 nan
nan nan nan nan nan 0.85985671
0.8565715 nan nan 0.8612861 0.86228581 nan
nan nan 0.86300008 0.86085716 nan 0.86242899
0.86271463 nan 0.86100016 nan nan 0.86157197
nan nan 0.86399997 nan nan nan
nan nan nan nan nan 0.86071477
0.86242857 nan nan nan 0.85771434 nan
0.86485724 0.86328565 nan 0.86171442]
warnings.warn(
Best parameters found by RandomizedSearch: {'n_estimators': 50, 'min_samples_split': 2, 'min_samples_leaf': 4, 'max_features': 'sqrt'
Confusion Matrix (Best Model):
[[2333  83]
 [ 316 268]]

Classification Report (Best Model):

```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.88 | 0.97 | 0.92 | 2416 |
| 1 | 0.76 | 0.46 | 0.57 | 584 |
| accuracy | | | 0.87 | 3000 |
| macro avg | 0.82 | 0.71 | 0.75 | 3000 |
| weighted avg | 0.86 | 0.87 | 0.85 | 3000 |

```

Accuracy Score (Best Model):
0.867

```

```

# Initialize SHAP explainer
explainer = shap.TreeExplainer(best_rf_model)

# Get SHAP values
shap_values = explainer.shap_values(X_test)

# To plot a binary classification SHAP we return a list of arrays

```

```
# If it is a binary classification, SHAP may return a list of arrays
# Select shap values for class 1 (index 1) for binary classification
if isinstance(shap_values, list):
    shap_values = shap_values[1]

# Check the shape of SHAP values and X_test
print(f"Shape of SHAP values: {shap_values.shape}")
print(f"Shape of X_test: {X_test.shape}")

# Ensure the shapes match
assert shap_values.shape[0] == X_test.shape[0], "Mismatch in the number of samples"
assert shap_values.shape[1] == X_test.shape[1], "Mismatch in the number of features"

# Summary plot for the selected class
shap.summary_plot(shap_values, X_test, feature_names=X.columns)
```

→ Shape of SHAP values: (3000, 11, 2)
Shape of X_test: (3000, 11)



```
# Save the model as a pickle file
import pickle
model_filename = 'best_rf_model.pkl'

# Open a file in write-binary mode and save the model
with open(model_filename, 'wb') as file:
    pickle.dump(best_rf_model, file)
```

